

ElasticSearch

Prof. Luca Cagliero
Dipartimento di Automatica e Informatica
Politecnico di Torino



**Politecnico
di Torino**

Lecture goal

- Information Retrieval
- Web search
- ElasticSearch

Information Retrieval

*Find material (usually documents) of an unstructured nature
(usually text) that satisfies an information need from
within large collections (usually stored on computers)*

What does “search” mean?



1. Take a query string
2. Match it against a document collection
 - Perform full-text search
 - handle synonyms
3. Calculate a set of relevant results
4. Score documents by relevance
5. Display a sorted list

Boolean retrieval

- Goal
 - Avoid scanning the text for each query
- Solution
 - Index the document in advance
- Method: Boolean retrieval model
 - pose any query in the form of a Boolean expression of terms
 - terms are combined with operators AND, OR and NOT
 - E.g., "Ceasar AND Brutus"
 - Returns all documents that satisfy the Boolean expression

Term-document incidence matrix

- Binary index representation
 - Avoid scanning the text for each query

	Documents →					
	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0

Term(s)
vector(s)

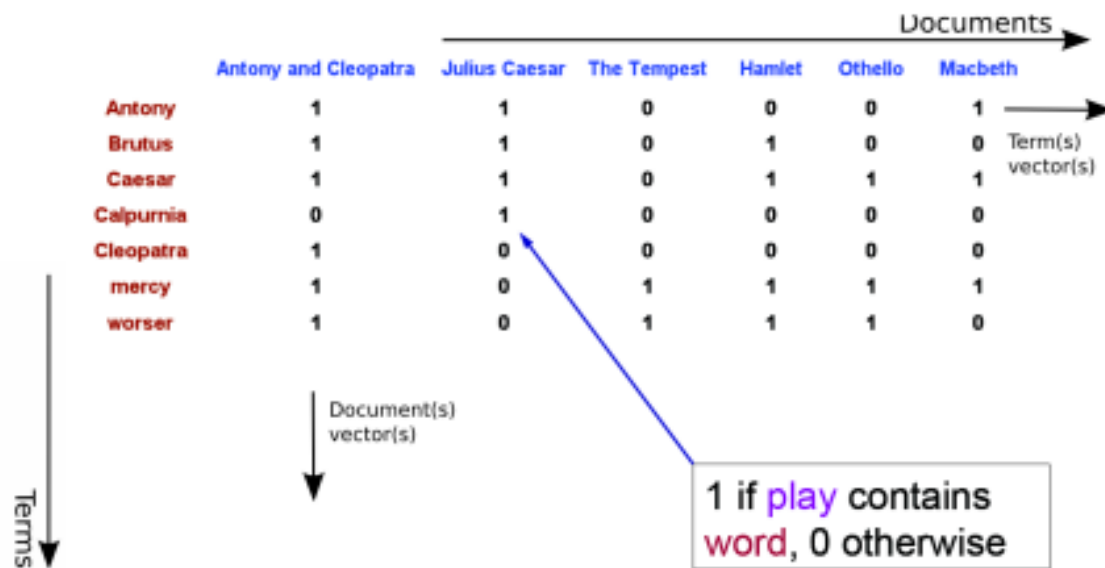
Document(s)
vector(s)

Terms

1 if play contains word, 0 otherwise

Term-document incidence matrix

- To answer the query *Brutus AND Ceasar AND NOT Calpurnia*:
 - Take the vectors for *Brutus*, *Ceasar* and the complement of the vector for *Calpurnia*: 110100 AND 110111 AND 101111 = 100100
 - The answer to this query are "Anthony and Cleopatra" and "Hamlet"



The diagram illustrates a term-document incidence matrix. A vertical arrow on the left points downwards and is labeled 'Terms'. A horizontal arrow at the top points to the right and is labeled 'Documents'. The matrix itself is a grid of 0s and 1s. The rows are labeled with terms: Antony, Brutus, Caesar, Calpurnia, Cleopatra, mercy, and worser. The columns are labeled with documents: Antony and Cleopatra, Julius Caesar, The Tempest, Hamlet, Othello, and Macbeth. A blue arrow points from a text box at the bottom right to the cell containing '1' in the row 'Calpurnia' and column 'Julius Caesar'. The text box contains the text: '1 if play contains word, 0 otherwise'.

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0

1 if play contains word, 0 otherwise

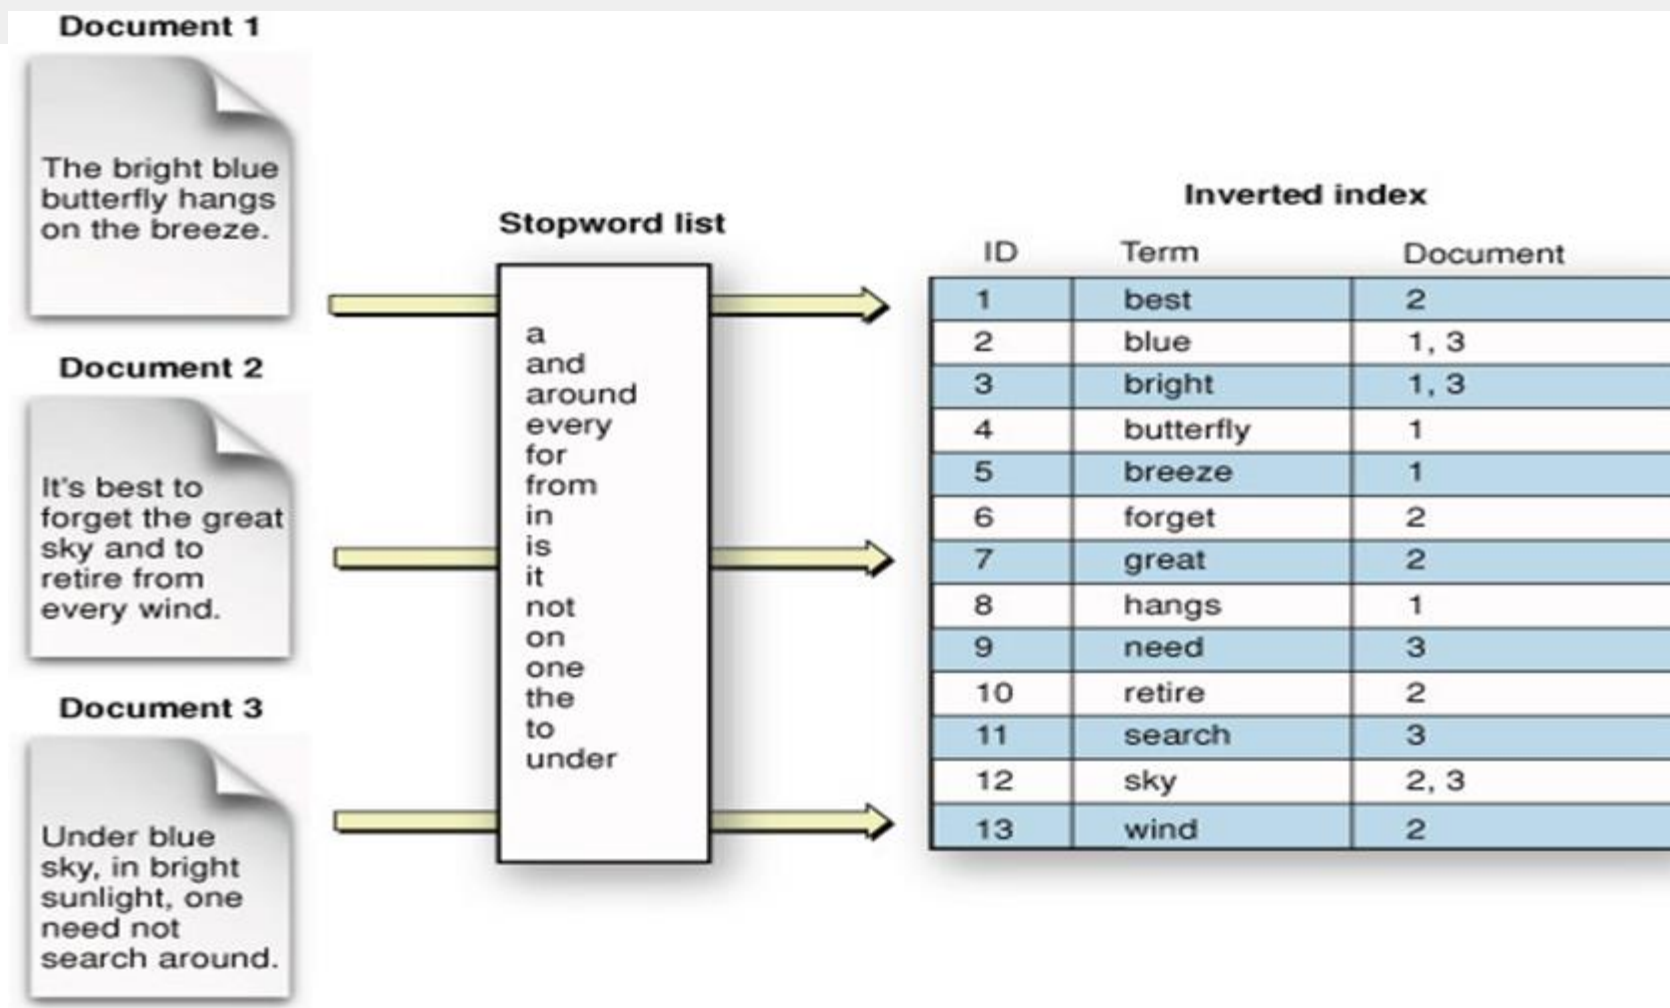
Boolean retrieval

- Drawbacks of the Boolean retrieval model
 - Sparse representation
 - Inefficient
 - All terms weigh equally
 - Capture syntactic text similarities only
- Solution to model sparseness
 - Inverted index
 - Stores only the 1s
- Solution to text similarities
 - Instead of 0s and 1s store the semantic similarities
 - e.g., BERT score

Boolean retrieval

- Solution to model sparseness
 - Inverted index
 - Stores only the 1s

Inverted index example



<https://stackoverflow.com/questions/47003336/elasticsearch-index-sharding-explanation> (latest access: April 2021)

Boolean retrieval

- Solution to term weighting
 - Use tf-idf weights or similar
- Solution to syntactic-only text similarities
 - Use semantic similarity measure
 - e.g., BERT score

Web search: challenges

- Content publishing on the Web is decentralized
- No central control of authorship
- Web page authors created content in dozens of (natural) languages and thousands of dialects

Web search

- The Web can be modelled as a directed graph in which each web page is a node and each hyperlink a directed edge
- Web search is commonly addressed using graph ranking algorithms
 - E.g., PageRank, HITS

More about graph ranking in the in-class practice



Additional reading on PageRank



- The Anatomy of a Large-Scale Hypertextual Web Search Engine. Sergey Brin, Lawrence Page. Computer Networks, vol. 30 (1998), pp. 107-117
- Download and read the paper: <https://snap.stanford.edu/class/cs224w-readings/Brin98Anatomy.pdf>

Additional reading on HITS



- Kleinberg, Jon (1999). "Authoritative sources in a hyperlinked environment" (PDF). *Journal of the ACM*. 46 (5): 604–632. doi:10.1145/324133.324140.
- Download and read the paper:
<http://www.cs.cornell.edu/home/kleinber/auth.pdf>

- Real-time distributed search and analytics engine
- Scalable and efficient data exploration
 - Full-text search
 - Highlighted search snippets
 - search as you type
 - did-you-mean
 - more-like-this
- Structured search
- Analytics
 - Real-time query answers on mixed data types
 - E.g., text, structured data

Elastic Search: popular example



- **GitHub** uses ElasticSearch to query 130+B lines of code.
- **Wikipedia** provides full-text search with highlighted snippets
- **StackOverflow** combines both full-text and geolocation queries for recommending related questions and answers



- Document-oriented (JSON) search engine
 - Complex data structures that may contain dates, geo locations, text, other objects, arrays of values
- Built on Lucene search engine library
 - Documents are indexed and searchable
- Highly available and horizontally scalable

- Data is stored in *named entries* belonging to a variety of data types
- SQL calls such an entry an *attribute* whereas in ElasticSearch it is called *field*
- In ElasticSearch a field can contain multiple values of the same type (list of values)
 - Similar to other NoSQL databases

ElasticSearch: field

SQL: column

- Data objects are represented as rows (SQL) or documents (ElasticSearch)
 - Row format in SQL is strict and follows a predefined schema
- Columns and fields are part of a row (SQL) or a document (ElasticSearch)
 - Documents are more flexible and can contain a variety of fields
 - They do not follow a strict schema

ElasticSearch: document

SQL: row

- An index is like a table in a relational database
- In Elasticsearch the indices are grouped in a cluster
 - As tables are included into a database

ElasticSearch: index **SQL:** table

ElasticSearch: cluster **SQL:** database

ElasticSearch: data representation



ES	cluster	index	document	field
SQL	database	table	row	column

- The index term has multiple meanings
 - (noun) An index stores a collection of documents
 - (verb) To index a document means to insert a document in an index
 - If the document already exists, it is replaced

- Additional structure that accelerates data retrieval
- Similar to a traditional relational index
- Every field in a document is indexed in ES
 - All inverted indices are used during search
 - Non-indexed fields (if any) are not searchable

- A document is the top-level (root) object serialized into JSON
- It is uniquely identified by the pair
 - Index: where the document (object) is stored
 - Id: the identifier of the document
 - Can be provided or uniquely generated by Elasticsearch

Three options:

- Structured query on specific fields
 - eventually sorted
 - similar to SQL query
- Full-text query
 - Finds all documents matching the search keywords
 - Returns them sorted by relevance
- A combination of the above

- Mapping
 - How the data in each field is interpreted
 - Elasticsearch dynamically generates a mapping by “guessing” data types
 - E.g., it may recognize a date type
- Analysis
 - How full text is processed to make it searchable
- Query DSL (Domain Specific Language)
 - Elastic Search query language


- Traditional data types (e.g., integer, float, date, but also string)
 - A value must **match exactly** the query
 - Similar to SQL
 - Examples: date, user ID, but also exact strings such as username or email address
- Question answered
 - “Does this document match the query?”

- Search for textual data
 - usually written in some human language
- Typical search is within the textual field
 - Examples: text of a tweet, body of an email
- It requires the definition of the concept of a document relevance to a query
- Question answered
 - How well does this document match the query?

- For solving full text queries it is very important to understand the underlying intent
 - abbreviations
 - e.g., USA vs United States of America
 - singulars/plurals, verb conjugation
 - e.g., cat vs cats, does vs did vs to do
 - synonyms
 - e.g., game vs competition
 - order of words building a context
 - e.g., fox news hunting vs fox hunting news

- Elasticsearch builds an inverted index on every full-text field
 - Designed for fast full-text search
- Inverted index
 - List of all the unique words that appear in any document in the collection
 - For each word it lists of the documents in which it appears

Full-text indexing



Two main steps:

1. Tokenization of a block of text into individual terms suitable for an inverted index
2. Normalization into a standard form to improve their retrieval (or recall) in queries
 - Terms are not exactly the same, but similar enough to be still relevant
 - Lowercase vs. uppercase
 - Stemming
 - e.g., cats vs cat
 - Synonym management
 - For searching, both indexed text and query string must be analyzed in the same way

Built-in functions provided by Elasticsearch:

- Character filter
 - clean the string before tokenization
- Tokenizer
 - split the string into individual words
 - E.g., by considering white spaces or punctuation as separators
- Token filters: operate on single terms
 - change terms (e.g., to lowercase)
 - remove (e.g., stopwords)
 - add terms (e.g., synonyms)

- A filter is used for fields containing exact values
 - It provides a boolean matches/does not match answer for every document
- A query is (typically) used for full-text search
 - It also asks the question: How well does this document match?
 - calculates how relevant each document is to the query
 - assigns it a relevance score, which is later used to sort matching documents by relevance
- The concept of relevance is well suited to full-text search
 - there is seldom a completely “correct” answer

- Filter execution is more efficient
- Filters are typically used to reduce the number of documents that have to be examined by a query
- Hint
 - use query clauses for full-text search or for any condition that should affect the relevance score
 - use filter clauses for everything else

- Expressed in Query DSL
- Submitted as formatted JSON in the body of an HTTP request

- Example: empty query
 - returns all documents in all indices

```
POST /_search
{}
```

- Search on a specific index

```
POST index1/_search
{}
```

- The top level field in an Elasticsearch query is always “query”
 - The query type is specified one level below
 - the query operates on the department index
 - specified in the URI
 - it performs the **search** operation

```
POST departments/_search
{
  "query": {
    "match" : { "name" : "John" }
  }
}
```

- Query
 - Find all the documents in the department index that have a field name containing the term John in it.
- Query type
 - Match query

```
POST departments/_search
{
  "query": {
    "match" : { "name" : "John" }
  }
}
```

- Complex queries specifying multiple matching criteria

```
POST departments/_search
{
  "query": {
    "bool": {
      "should": [
        {"match": {"name": "John"}},
        {"match": {"name": "Mark"}}
      ],
      "minimum_should_match": 1,
      "must": {
        {"match": {"title": "developer"}}
      },
      "must_not": {
        {"match": {"lastname": "Smith"}}
      }
    }
  }
}
```

bool specifies the compound query

Compound queries



POST departments/_search

```
{
  "query": {
    "bool": {
      "should": [
        {"match": {"name": "John"}},
        {"match": {"name": "Mark"}}
      ],
      "minimum_should_match": 1,
      "must": {
        {"match": {"title": "developer"}}
      },
      "must_not": {
        {"match": {"lastname": "Smith"}}
      }
    }
  }
}
```

should specifies the OR condition

must corresponds to the AND condition

must_not specifies the NOT condition

- Can be used for both full-text and exact queries
- On a full-text field
 - it analyzes the query string with the correct analyzer before executing the search
 - it returns a relevance score `_score` for the search
- On an exact field or a not analyzed string field
 - it searches the exact value
 - it returns a relevance score `_score` of 1
- When a bool query is specified on full-text fields
 - It combines the `_score` from each must or should clause that matches

- It is possible to specify multiple indices to be searched in the query URI

```
POST rooms,students/_search {...}
```

- When a number of documents can be returned as query result, by default the top 10 relevant results are returned
- Earlier versions of Elasticsearch include index types that have been deprecated since version 7.0

- Insert of a new single document is performed by means of a POST operation
 - Name of index
 - JSON document to be indexed
- **Index_name**: name of the index in which the document should be inserted
- **<id>**: optional parameter that associate the document with a specific identifier
 - If the ID is not provided, Elasticsearch creates a unique identifier for the document (e.g., W0tpsmlBdwcYyG50zbta)

- Documents in ES are immutable
 - To update a document, it is reindexed
- When a document is updated, ES
 - Retrieves the old document
 - Modifies the retrieved copy
 - Deletes the old document
 - Indexes the new document (the copy)
- Internally, the old version of the document is not deleted immediately
 - It is not accessible
 - Deleted documents are cleaned in background

- The update of a document is performed using a POST request
 - Name of the index
 - Unique ID of the document
 - the fields to be updated and the associated new values
 - To update a document, it is reindexed

```
PUT index_name/123/_update
{
  "color" : "red",
}
```

This update request modifies the document with ID=123 by setting the value of the “color” field to “red”


- The deletion of a document is performed using a DELETE request
 - Name of the index
 - Unique ID of the document

```
DELETE index_name/id
```

- The operation removes a JSON document from the specified index
 - Document removal is not immediate!

- In Elasticsearch the relevance score is a floating-point number
 - computed for each document matching the query
 - stored as `_score` for each document in the search result
 - higher `_score` values correspond to more relevant documents
- Sorting by relevance is performed by considering the `_score` variable
 - by default, documents in a query result are sorted by descending value of the `_score` field

1. Compute matching results for the query
 - Compute relevance score for all documents in the query result
2. Select top relevance documents (hits)
 - Default is 10 hits (documents)
3. (optional) Re-score documents
 - more computationally expensive algorithm



- Need to compute the similarity between
 - The query
 - Each document
- Each document may contain a (different) subset of the query terms

1. Select documents matching the query
 - Boolean model
 - Fast computation
2. Evaluate the importance (weight) of each term in a document with respect to the query
 - Term importance evaluated using the TF/IDF (Term Frequency/Inverse Document Frequency) score
 - Document and query are represented in vector form
 - Vector Space Model
3. Evaluate the similarity of the vector representation of the query and the document


- Key factors
 1. Term frequency
 2. Inverse document frequency
- They are calculated and stored at index time
- They are used to calculate the weight of a single term in a document
 - Other methods can be used
- Queries usually contain more than one term
 - Need a way to combine multiple terms

- It represents both query and document as (term) vectors
- It provides a way to compare a multi-term query against a document
- A query (or document) is represented as a vector
 - The vector size is the number of terms in the query
 - Each vector element is the weight of one term, calculated with BM25 scoring
 - Extensions towards semantic similarity (e.g., BERT similarity) are possible!
- Vectors can be compared by measuring the angle between them
 - Cosine similarity


- Vectors can be compared by measuring the angle between them
 - Cosine similarity
- The angle between a document vector and a query vector is used to compute the similarity between a document and a query
 - It assigns to the document its relevance score for the query

Vector Space Model

- Query example
 - *Happy hippopotamus*
- Document examples
 - *I am happy in summer.*
 - *After Christmas I'm a hippopotamus.*
 - *The happy hippopotamus helped Harry.*
- It is possible to create a vector for each document
 - Document 1: (happy, _____) \rightarrow $[BM25_{happy}, 0]$
 - Document 2: (__ , hippopotamus) \rightarrow $[0, BM25_{hippopotamus}]$
 - Document 3: (happy, hippopotamus) \rightarrow $[BM25_{happy}, BM25_{hippopotamus}]$




Horizontal scalability



- Sharding is a technique to divide an index in smaller partitions
 - Each partition is a shard
- Each document belongs to a single shard
 - Each shard is an instance of a Lucene index
- When data is written to a shard
 - It is periodically (every 1 second) written into a new immutable Lucene segment on disk
 - It becomes available for querying
- Shards are the elementary units in which data is distributed on nodes in a cluster

- A cluster is a collection of multiple machines
 - i.e., nodes in the cluster
- Shards can be stored in any node within the cluster



Why sharding?



- It allows splitting data in smaller chunks, and thereby scaling on large volumes of data
 - Data may be distributed across multiple nodes within a cluster
 - Shards can be stored on smaller disks
 - E.g., it is possible to store 1TB of data even without a single node with that disk capacity
- Operations can be distributed across multiple nodes and thereby parallelized
 - Performance is increased, because multiple machines can potentially work on the same query
- Shards may be replicated on different nodes to increase availability

- Elasticsearch uses optimistic concurrency control
 - It assumes that conflicts are unlikely to happen
 - However, if the underlying data has been modified between reading and writing, the update will fail
- Different from ACID transactions that need locking
- The process is “simple” for centralized data management

- Elasticsearch data may be distributed on different nodes in a cluster
 - Shards may be replicated on different nodes (replica shards)
- When documents are created, updated, or deleted, the new version of the document has to be replicated to other nodes in the cluster
 - The primary copy is always written first
 - The replication requests are sent in parallel and may arrive at their destination out of sequence

Document versioning



- Elasticsearch needs a way of ensuring that an older version of a document never overwrites a newer version
 - Every document has a `_version` number that is incremented whenever a document is changed
- Elasticsearch uses this `_version` number to ensure that changes are applied in the correct order
 - if an older version of a document arrives after a new version, it can be ignored
 - the `_version` number is used to ensure that conflicting changes made by applications do not result in data loss
- APIs that update or delete a document accept a version parameter
 - can apply optimistic concurrency control only when needed

Acknowledgements and copyright license

- Copyright licence

- Attribution + Noncommercial + NoDerivatives



- Acknowledgements

- I would like to thank Dr. Moreno La Quatra, who collaborated to the writing and revision of the teaching content

- Affiliation

- The author and his staff are currently members of the Database and Data Mining Group at Dipartimento di Automatica e Informatica (Politecnico di Torino) and of the SmartData interdepartmental centre
 - <https://dbdmg.polito.it>
 - <https://smartdata.polito.it>

Thank you!